

Package Hallucinations Revisited: Measurement Validity, Replication, and Tail Risk in Modern Code Models

Jacob Anderson

jwa@beyond-ordinary.com

Manuscript draft, September 2026

Abstract

When a code model recommends a package that does not exist on the registry, anyone can register that name and fill it with malicious code. Several studies now report how often models do this. We show that the reported rate depends heavily on how it is measured. We re-ran the 2024 study by Spracklen et al. under its own settings, with three random seeds per model. We built a stricter parser that refuses to mine package names out of prose or code, and we checked it against labeled samples. We ran a controlled test of the token cap that earlier work had blamed for large rate changes. We then measured six current models on the same 800 Python prompts, with two package questions per prompt, under a plan fixed before any data were collected. Three later cells added Qwen 3.5, Kimi K2.7, and Kimi K3. The old results mostly hold up. CodeLlama came out at 24.55% against a published 26.12%. DeepSeek kept its rank, but under its own original parser it scored 13.28–13.54%, below the published 16.61%, and every seed’s confidence interval excluded that value. The token-cap effect that we and an independent audit had reported earlier [1, 17] did not reappear as a consistent or monotone effect. Modern models ranged from 1.31% to 17.91%, but the average hides three different kinds of failure. For Grok 4.6 and DeepSeek V4 Flash, three answers supplied 82.88% and 72.42% of all unregistered names; DeepSeek Coder V2’s 10.30% was spread across many prompts. Kimi K3 pushed the range to 33.30%, yet 46 answers that hit the token cap supplied 81.02% of its unregistered names, and it put fewer prompts at risk than its predecessor, so the ranking between the two flips with the choice of metric. Rescoring the same outputs with the extractor of a concurrent 2026 replication kept the original six cells at 8.30–17.49%, and four of the five names that all six models produced came from one identical prompt. A single average rate is not enough. Rates need a dated registry, a stated parser with its coverage, confidence intervals built on prompts, and a measure of how concentrated the errors are.

1 Introduction

1.1 Motivation

Code models do more than write code. Ask one which packages you need, and it answers with names. If one of those names is not registered on the package index, anyone can register it. This is the *slopsquatting* threat [6, 13]: an attacker publishes a malicious package under a name that models are already telling developers to install [10, 11, 18]. Typosquatting [9, 14, 16, 20] waits for a developer to mistype a name. Slopsquatting does not have to wait. The model supplies both the name and the confidence.

Two distinctions matter before we count anything. First, a recommendation is not the same as an import. An install command, or an answer to “which packages do I need,” is an explicit recommendation. An import statement inside generated code might just refer to a file in the same project. Second, a name that is not on PyPI is not necessarily made up. Models often name real tools from other ecosystems, such as `apt`, `npm`, Java, or Go packages, that simply cannot be installed with `pip`. They also name Python standard-library modules as if they were packages. All of these are unsafe recommendations, but they are different mistakes. So we measure an *unregistered-PyPI recommendation rate*, and we use the word “hallucination” only for the older metric that we replicate.

This is also why replication is hard. A reported rate depends on four choices: how the model’s free text is turned into package names, which snapshot of the registry counts as the truth, how many tokens the model was allowed to produce, and how the errors are spread across answers. None of these is a small detail. Each of them changed a conclusion in the work reported here.

1.2 Gap

Spracklen et al. [18] established the problem at scale. Across 16 models and 576,000 code samples, 19.7% of recommended packages did not exist, and 205,474 different invented names appeared. A 2026 replication on five frontier models [4] reports that the gap between models has narrowed and that some names are invented by every model. Both

studies report one average rate per model. Our own first attempt to replicate the 2024 result through a plain HTTP client [1], checked by an independent audit [17], found four ways that an average can mislead:

- A permissive parser that splits text on commas counts code fragments and prose as invented packages.
- A token cap can silently produce empty answers from a reasoning model, or cut a list in the middle of a name.
- The registry changes over time, so a name can switch from “invented” to “real” with no change in the model.
- A 2024 checkpoint served through a 2026 stack is not provably the same model.

An average also cannot tell you whether the risk is spread evenly or concentrated in a few runaway answers. Those two situations call for different defenses.

1.3 Contributions

1. A properly powered replication of the 2024 result, with three seeds per model, that separates “the pattern holds” from “the exact number holds” (§4.2).
2. A strict, validated parser that refuses unclear answers instead of guessing, checked again on the newer hosted-model output formats after the campaign (§3.4, §4.1).
3. A controlled experiment on the token cap, with identical prompts and deterministic decoding (§4.3).
4. A benchmark of six current models on the same prompts, scored against two registry snapshots, with confidence intervals that respect the prompt structure and a way to tell diffuse failures from catastrophic ones (§4.4–§4.6). Three later cells extend it, and the comparison between Kimi K2.7 and Kimi K3 shows a ranking that flips with the choice of metric (§4.9, §4.10).
5. A direct comparison with the concurrent Churilov study, made by rescoring our own outputs with that study’s extractor, plus an overlap analysis that asks whether names shared across models come from shared prompts (§4.7).

1.4 Result preview

- CodeLlama 7B: 24.55% (mean of three seeds) against a published 26.12%.
- DeepSeek Coder 6.7B under its original parser: 13.43% against a published 16.61%. All three seed intervals exclude the published value.
- Token cap: no consistent or monotone effect in the controlled test.
- Modern models: primary rates from 1.31% to 17.91%, with very different concentration.
- Rescoring with Churilov’s extractor: the original six Campaign 4 cells stay at 8.30–17.49%. The extension cells score 6.51% (Qwen 3.5), 7.38% (Kimi K2.7), and 8.37% (Kimi K3). Of the 621 flagged mentions, 608 (97.91%) come from imports rather than install commands.

- Overlap: five names appear in all six Campaign 4 models, and four of them appear in all six on the same prompt.
- Cloud extension: Qwen 3.5 at 2.77% (95% interval 2.07–3.54). Kimi K2.7 at 23.10%, but only 4.95% among answers with at most ten packages, and rising steeply with list position.
- Kimi K3: 33.30% (27.22–38.81) when every recommended name counts equally, 21.50% of prompts at risk, and a 4.18% mean rate per answer. Setting aside its 46 cap-hit answers, as a diagnostic only, gives 13.11%, next to K2.7’s 13.18%.

2 Background and related work

2.1 Package hallucinations and supply-chain risk

The original study [18] counts package hallucinations in three ways. It parses `pip install` and `npm install` commands out of generated code. It asks the model which packages its own generated code requires. And it asks the model which packages would be useful for the original task. Every extracted name is normalized and looked up in a master list of registry names taken on 2024-01-10. Any name that is missing from the list counts as a hallucination. The threat model is an attacker with no special access who can register packages. The study withheld the hallucinated names themselves for safety. We adopt both the threat model and the release policy: no candidate unregistered name appears in this paper or its artifacts. The original study describes the threat as a variation of the classical package confusion attack. The name *slopsquatting*, modeled on typosquatting, was coined by Seth Larson and popularized by Andrew Nesbitt in April 2025 [6, 13], roughly ten months after the study’s preprint appeared.

Three terms recur below. A *package recommendation* is a name the model offered when asked a package question. A *generated import* is a name that appears in an import statement of generated code. An *unregistered name* is absent from the relevant PyPI snapshot, and an *ecosystem mismatch* is an unregistered name that is a real package somewhere else. All of these matter for slopsquatting, but only rates built on the same channel can be compared across studies.

2.2 Evaluation of code-generating models

Code benchmarks [3] pin a model version and report one number per model. Hosted models make pinning harder, because the same alias can serve a different model next month. Decoding is random, so one run is one draw, and two models should be compared on the same prompts, pair by pair. Most important for this measurement, the package names in one answer are not independent observations. They share a prompt, a task, and a single generation. Treating each

name as its own coin flip is pseudoreplication [8]. The right unit for resampling is the cluster, here the prompt [2, 5].

2.3 Measurement error in free-form generations

Turning free text into package names is a classification task, and it has its own precision and recall. A permissive parser catches almost everything but invents false positives whenever the model stops following the requested format. A strict parser refuses answers that do not follow the format and misses real recommendations buried in prose. Either way, the choice is part of the measurement and belongs next to the rate. Registry membership also changes over time. Names get registered and deleted after any snapshot, so a name can move from “hallucinated” to real, or from real to deleted, while the model stays the same. Finally, the shape of the response matters. A token cap cuts lists short. A reasoning model can spend its whole budget on hidden reasoning and return nothing. A hosted endpoint can leak a literal control token into the text. Each of these looks like model behavior if nobody checks.

2.4 Concurrent 2026 replication

Churilov [4] re-ran the full original prompt corpus (199,845 Python and JavaScript generations) on five frontier models released between October 2025 and March 2026. That study extracts package references from generated code (install commands plus imports) and checks them against the original master lists. Beyond the replication, it finds 127 names that all five models invent, of which 53 could still be registered after coordinated disclosure with PyPI Security and Socket.dev. It also reports that Python now has a higher rate than JavaScript, and that Claude Haiku 4.5 scores below Claude Sonnet 4.6. It is the broader study in languages, corpus size, registrability testing, and disclosure, and it shows that the cross-model attack surface is real.

Three points matter for the comparison in §4.7. First, its code-import measure is not the original study’s three-channel measure, and it is not our package-recommendation measure either. A rate from one is not a rate from the other. Second, its headline that the range between models has compressed (4.62–6.10% overall, 5.49–7.27% on Python, against 5.2% and 21.7%) compares five frontier models with the original study’s *category averages* for commercial and open-source models. The original model-level range on Python was 3.59% to 26.12%. Third, its Wilson intervals, chi-square tests, and permutation tests treat every extracted mention as an independent observation, even though mentions cluster inside generations. We do not claim its significant pairs would disappear under cluster-aware inference, only that they were not computed that way. Its reading of name overlap as evidence of shared training data is one hypothesis. Shared

prompts, extractor behavior, and corpus composition are alternatives that §4.8 tests directly.

2.5 Research questions

RQ1 (Replication). Do the published CodeLlama and DeepSeek results come back when we use a larger fixed prompt set and repeat the generation with several seeds?

RQ2 (Measurement validity). How much do the parser, its strictness, the response format, and the registry date change the reported rate?

RQ3 (Cap sensitivity). Does the rate change with the token cap on package answers, once the code and prompts are held fixed and decoding is deterministic?

RQ4 (Modern models). How do current hosted and local models compare on one prompt set, one frozen parser, and one two-registry definition?

RQ5 (Failure shape). Are the rates driven by errors spread across many answers, or by a few catastrophic recommendation floods?

RQ6 (Concurrent-work bridge). When we rescore identical outputs with Churilov’s extractor, and ask whether shared names came from shared prompts, which agreements and disagreements remain?

RQ7 (Successor and list-volume mechanism). Do Kimi K2.7’s list-flood and late-position effects persist in Kimi K3, once we record the finish reason and token count for every answer?

3 Study design

3.1 Two-track design

Figure 1 shows the design. We wrote it down and froze it before collecting data, and we logged every later change as a dated amendment; this is the practice known as preregistration [15]. **Track A** is the historical replication. It uses the original model families, the original sampling settings and token caps, the original parsers, and the frozen 2024 registry. It answers one question: does the 2024 result come back? **Track B** is the modern benchmark. It uses the same prompts, our validated parser (Parser v2), larger token budgets, and both the 2024 and the 2026 registries. It allows comparisons between current models under one instrument. It is not, and does not claim to be, a reproduction of the 2024 study at its original scale. The **extension cells** came later, each under its own frozen protocol. Qwen 3.5 Cloud and Kimi K2.7 Code Cloud ran under a protocol frozen on 2026-08-14. (That protocol first named Kimi K3, but the K3 endpoint returned HTTP 402 before generating anything, so we substituted K2.7 and logged the change.) Kimi K3 ran afterwards under a separate specification written to study the K2.7 mechanism in its successor. These cells extend Track B; they do not




Figure 1: The two-track design and the measurement pipeline. Every cell uses the same fixed set of 800 Python prompts and asks two package questions per prompt. The tracks differ in token caps, parser, registries, and in what they allow us to claim. Track B is not a scale reproduction of the 2024 study.

change its original set of comparisons.

For Kimi K3 we use the word “frozen” carefully. The scientific design was fixed before we looked at any outcome: prompts, parser, registries, model and sampling settings, the quantities to estimate, the diagnostics, and the set of comparisons. But while collection was running, and before any package answer was scored, we expanded the code that records provenance and checks completeness. The running collector had already loaded the earlier version. Those edits changed no prompt, response, parser decision, or outcome calculation, but the executable pipeline was not frozen byte for byte. And because the K3 preregistration was not committed to version control before collection, the run is preregistered in spirit, not cryptographically. A provenance stamp written afterwards hashes the attested design, the local timeline, every file in the commit, and a content-addressed index of the 93 raw-artifact files. It says plainly that it cannot repair the missing pre-data commit. The repository now enforces a commit-based freeze gate for any future analytic run.

3.2 Prompt corpus and unit of analysis

The original artifact provides four Python prompt datasets: Stack Overflow questions (all-time and last-year popularity) and LLM-generated prompts based on popular packages over the same two periods. The Stack Overflow datasets are sorted by popularity, so taking the first n prompts would bias the sample. Instead we draw 200 prompts per dataset with a fixed random seed and use that one subset in every cell, seed, and phase. Each prompt yields one code generation and two package answers, so a complete cell has 800 prompts and 1,600 package answers. Every confidence interval and every comparison treats the prompt as the unit.

3.3 Models and identity capture

Table 1 lists every cell: the exact model identity served, its size and quantization where known, its reasoning setting, the sampling parameters the provider accepted, the token caps, and the date. Track A uses CodeLlama 7B Instruct and DeepSeek Coder 6.7B Instruct, served locally by Ollama, with the immutable digest, quantization, and server version

recorded. Track B uses Claude Opus 5, GPT-5.2, gpt-oss 20B, Grok 4.6, DeepSeek Coder V2 16B, and DeepSeek V4 Flash. For hosted models we record the model id that the provider echoes in every response. We turned reasoning off wherever the provider offered a switch (Anthropic thinking disabled; OpenAI reasoning effort none; `think:false` on Ollama Cloud). gpt-oss kept its native reasoning behavior, so reasoning was not disabled everywhere. When a provider rejected a sampling parameter, the runner dropped it, retried, and recorded the change. The DeepSeek Coder V2 cell was resumed across four invocations with different worker counts and context settings. Its model digest never changed, and we disclose the mixed settings rather than hide them.

3.4 Parser v2 and response statuses

Parser v2 accepts an answer only if it looks like a list of package names, separated by commas or line breaks, where every item is a valid name in the PEP 503 sense [19]. List markers such as “12.” or “-” are removed only at the start of an item. That fixes an old normalizer that turned “12. requests” into “1requests.” Quotes and brackets are removed only in matched pairs. Names are counted once per answer. Every answer gets one of three labels: *list*, *malformed* (code fences, prose, or any item that is not a name), or *empty*. A malformed answer contributes nothing to the rate; we report how often it happens instead. This is a deliberate trade: we give up some recall to avoid inventing false positives, and §4.1 measures the cost.

3.5 Registries and categories

We use two PyPI snapshots: the frozen list from the original artifact, taken on 2024-01-10 (500,513 names), and a snapshot taken at the start of our campaign on 2026-08-12 (869,894 names). The namespace grew by 74% in between. Our primary metric counts a recommended name only if it is missing from *both* snapshots and is not a Python standard-library module. Each condition earns its place. Since 2024, 8,961 names on the frozen list have been deleted from PyPI. Some of those were squatted registrations of standard-library names such as `os` and `json`. That means the original method

Cell	Track	Identity	Size, quant.	Reasoning setting	Sampling sent	Caps	Date
CodeLlama 7B Instruct	A	codellama:7b-instruct; 8fdf8f75...	digest 7B, Q4_0 GGUF	n/a	all three	2048 / 64	2026-08-13
DeepSeek 6.7B Instruct	A	deepseek-coder:6.7b-instruct; digest ce298d98...	7B, Q4_0 GGUF	n/a	all three	2048 / 64	2026-08-13
Claude Opus 5	B	claude-opus-5 (Anthropic API)	hosted	thinking disabled	T, top- <i>k</i> , top- <i>p</i> rejected, dropped	4096 / 2048	2026-08-13
GPT-5.2	B	gpt-5.2-2025-12-11 (OpenAI API)	hosted	reasoning effort none	top- <i>k</i> dropped	4096 / 2048	2026-08-13
gpt-oss 20B	B	gpt-oss:20b; digest 17052f91...	20.9B, MXFP4 GGUF	native reasoning retained default	all three	4096 / 2048	2026-08-13
Grok 4.6	B	grok-4.6 (xAI API)	hosted		top- <i>k</i> dropped	4096 / 2048	2026-08-13
DeepSeek 16B	B	deepseek-coder-v2:16b; 63fb193b...	digest 15.7B, Q4_0 GGUF	n/a	all three	4096 / 2048	2026-08-13 [†]
DeepSeek V4 Flash	B	deepseek-v4-flash:cloud (Ollama Cloud)	hosted	think:false	all three	4096 / 2048	2026-08-13
Qwen 3.5 Cloud	C	qwen3.5:cloud, served qwen3.5	397B, BF16	think:false	all three	4096 / 2048	2026-08-19
Kimi K2.7 Code	C	kimi-k2.7-code:cloud, served kimi-k2.7-code	1.042T, INT4	think:false	all three	4096 / 2048	2026-08-19
Kimi K3	D	kimi-k3:cloud, served kimi-k3	2.812T, MXFP4	think:false	all three	4096 / 2048	2026-08-25

Tracks: A = historical replication (paper protocol, three generation seeds); B = Campaign 4 modern benchmark; C = separately preregistered Ollama Cloud extension; D = Kimi K3 successor cell. Sampling: temperature 0.7 (code) / 0.01 (package queries), top-*k* 20, top-*p* 0.9 where the provider accepts them. Caps are code / package-query response tokens. Every hosted cell records the served model id echoed in each response; local cells record the Ollama digest, quantization, and server version. [†]Four resumed invocations with mixed worker counts (1, 2) and context settings (12,288; 16,384; 32,768); the digest was stable throughout.

Table 1: Model cells, exact identities, and protocol settings.

silently counted standard-library confusion as valid packages, because its Python pipeline, unlike its JavaScript one, did not exclude core modules. We also record secondary categories: valid in both snapshots, registered after the 2024 snapshot, deleted since the snapshot, standard-library confusion, and unregistered in both.

3.6 Outcomes

Every cell reports the same set of numbers. The *occurrence rate* is the share of accepted package names that are unregistered; an answer with 600 names counts 600 times. *Prompt risk* is the share of prompts with at least one unregistered name. We report the share of malformed and empty answers, the number of packages per prompt and per accepted answer, and the number of distinct unregistered names. To describe concentration we report the share of unregistered names that come from the worst one, three, and ten answers, and the rate that remains after removing them (*leave-top-k*). For the cap diagnostic and the extension cells we also report rates by list position. Leave-top-*k* numbers describe concentration; they are never a substitute for the primary rate.

3.7 Statistical analysis

Every confidence interval is a bootstrap that resamples prompts, not package names. We resample prompts with replacement inside each of the four 200-prompt datasets, recompute the quantity of interest, and repeat (50,000 times for cell intervals and pairwise comparisons; 20,000 for the cap diagnostic). When we compare two cells, we pair them prompt by prompt on the 800 shared prompts and apply Holm’s correction [7] across the 15 Track B pairs, because

testing many pairs at once inflates the chance of a false positive. The preregistration fixed the bootstrap but not the exact rule for turning it into a *p*-value, so we report both the preregistered two-sided bootstrap-tail probability and, clearly labeled as a check done afterwards, a recentered-null bootstrap with the same replicates, strata, pairing, and seed. Effect sizes and intervals come first. When an adjusted *p*-value hits the floor of what 50,000 replicates can resolve, we report it as a bound: for the eight-comparison K3 family that floor is $8 \times 1/50,000 = 0.000160$, written $p \leq 0.000160$, never as an exact value. The K2.7-versus-K3 comparisons of prompt risk and flood frequency use the same paired, dataset-stratified resampling.

3.8 Churilov bridge and prompt-conditioned overlap

To compare with the concurrent study without generating anything new, we took the frozen generated code from every cell and rescored it with a port of Churilov’s public extractor: a `pip install regex` plus top-level imports (parsed with the Python AST, falling back to tree-sitter), the same standard-library list and false-positive filter, imports counted once per generation, and no deduplication between the install and import channels. We report two rates: absence from the frozen registry, which is the closest match to their artifact, and our own two-registry, non-stdlib rate. Both come with prompt-based intervals, stratified by dataset. For overlap, we report the ordinary Jaccard similarity between the sets of unregistered names of each model pair, and beside it an aligned Jaccard over (normalized name, dataset, prompt index) tuples. For every pair, and for the names shared by all models, we count how many shared names occurred on at

least one identical prompt. We publish no candidate names.

3.9 Cloud extension and per-response provenance

The Qwen 3.5 and Kimi K2.7 cells reuse the fixed prompt subset, Parser v2, both registries, and the two-question design. Comparisons that involve a new cell form their own family of 13 Holm-corrected tests. We report Query 1 and Query 2 volumes separately, because once a model lists hundreds of names in one channel, the two questions stop behaving like repeat measurements. Kimi K2.7’s combined preregistered estimate stays as it was; its breakdowns by question, list length, and position were done afterwards. The first extension runner recorded cap hits only per phase, not per answer, so for K2.7 we can only *approximate* which answers hit the cap by taking the longest answers in each phase. Kimi K3 fixed that. Every K3 answer carries a sidecar record with its finish reason, served model, prompt and completion tokens, and exact cap status, and scoring refuses to run if any of the 2,400 records is missing. We also recorded the provider’s token counts against a dollar budget the maintainer had authorized in advance.

4 Results

4.1 Measurement validation

The instrument changed the conclusions, so we describe it first (Table 2). On the initial validation sample of 200 answers, Parser v2 assigned the right status 98% of the time. All four mistakes had the same cause: quote characters wrapped around an entire list. No answer labeled as a list contained a wrongly extracted name. Among names that appeared in a valid list format, the parser found 97.78%. Counting every recommendation a human could see, including those buried in prose or code, it found 79.75%: 102 recommendations sat inside text that the grammar refuses to mine. That is the price of strictness, and it is why we report malformed rates next to every rate below. The validation sample was drawn before the hosted models were run, so after the campaign we audited 180 more answers from Track B (20 chosen at random and 10 deliberately difficult ones per model). All 180 matched the frozen parser on both status and the exact names extracted.

The audit also found two problems that agreement with the parser cannot reveal. Grok 4.6 ended 184 answers with a literal `<|eos|>` marker. All 184 count as malformed under the frozen grammar, and all 184 become valid lists when that marker, and only that marker, is removed. Removing it drops Grok’s malformed rate from 12.56% to 1.06%, but moves its primary rate only from 8.22% to 7.91%, because the 262 names recovered are registered or standard-library names. DeepSeek Coder V2 wrote the collapsed

Quantity	Value
<i>Initial validation sample (pre-campaign runs; labels signed 2026-08-13)</i>	
Records	200
Response-status accuracy	98.0% (196/200)
Status errors (root cause)	4 (wrappers around an entire comma list)
Extraction precision	100% (no incorrectly extracted token)
Grammar-valid occurrence recall	97.78% (441/451)
End-to-end semantic recall	79.75%
after the wrapper fix (v2.1 candidate)	81.56%
Occurrences withheld by deliberate abstention	102
Records with recommendations inside malformed text	42
<i>Post-campaign Track B audit (hosted formats; 20 random + 10 challenge per cell)</i>	
Records	180
Frozen-contract status agreement	100% (180/180)
Exact extraction agreement	100%
Random core / challenge set	120/120 / 60/60
Semantic anomalies annotated	4 no-package prose; 4 literal <code>< eos ></code> markers; 1 collapsed no-package sentinel

Both reviews were delegated to a second AI system (OpenAI Codex) at maintainer direction; no human read every record. Contract agreement validates the frozen grammar, not semantic validity of every extracted token.

Table 2: Parser v2 validation and the Track B frozen-contract audit.

token `NoPythonpackagesrequired` eleven times, which the grammar accepts as a name. Treating such one-word answers as “no packages” moves its rate from 10.30% to 9.89% (and gpt-oss from 2.15% to 2.14%). The frozen results stay primary; we report both adjustments as sensitivities.

The registry date changes what counts as an error (Table 3). Standard-library names make up 10–20% of accepted names in every cell, most of all in DeepSeek Coder V2, and the original method could not see them. Scoring against the 2026 snapshot alone inflates every rate, because the deleted squats of standard-library names now look like missing packages. Scoring against the 2024 snapshot alone counts packages registered since then as hallucinations and counts squatted standard-library names as valid. Only the two-snapshot, non-stdlib definition is stable in both directions.

4.2 Historical replication

The 2024 pattern comes back; the exact DeepSeek number does not (Figure 2 and Table 4). Under the paper’s exact protocol, with three independent seeds on the fixed prompt set, CodeLlama 7B scored 24.20%, 24.84%, and 24.60% (mean 24.55%). Every one of its intervals contains the published 26.12%. DeepSeek Coder 6.7B scored 14.37–14.86% under the generic parser. Under its own original parser, which is the fair test, it scored 13.28–13.54%, and all three intervals (roughly 11.6–15.4%) exclude the published 16.61%. The seeds differ from each other by only 0.5–0.6 percentage points for both models, so regeneration noise is small, and CodeLlama scores above DeepSeek at every seed.

So the ordering and the rough size of the effect reproduce, but the exact DeepSeek value does not. Both models come

Model	Occurrences	Valid both (%)	Stdlib (%)	Post-snap.	Deleted	Unreg. both	Frozen-only (%)	Current-only (%)	Primary (%)
Claude Opus 5	9,280	83.5	14.8	15	19	122	1.64	10.77	1.31
GPT-5.2	7,222	87.9	9.9	10	7	136	2.12	6.63	1.88
gpt-oss 20B	6,223	80.2	17.4	6	12	134	2.41	11.60	2.15
Grok 4.6	6,754	77.6	13.7	28	5	555	8.75	16.42	8.22
DeepSeek Coder V2 16B	2,437	69.1	19.8	9	12	251	10.79	24.25	10.30
DeepSeek V4 Flash	6,054	68.9	12.8	14	14	1,084	18.20	25.98	17.91

Categories partition all Parser v2-accepted occurrences. Stdlib = standard-library module named as an installable package. Post-snapshot = registered on PyPI after 2024-01-10; deleted = present in the frozen list but absent from the 2026-08-12 snapshot. Frozen-only and current-only rates score absence from a single snapshot with no stdlib exclusion; the current-only column is inflated because squat-registrations of stdlib names present in 2024 have since been purged. The primary rate requires absence from both snapshots and non-stdlib status.

Table 3: Name-category composition by model and the effect of registry choice.

Model	Paper (%)	Seed	Generic (%)	95% CI	⊃ paper	Family (%)	95% CI	⊃ paper	v2 malformed (%)	Cap hits
CodeLlama 7B	26.12	101	24.20	22.09–26.45	yes	–	–	–	10.9	54
		102	24.84	22.55–27.26	yes	–	–	–	8.5	60
		103	24.60	22.23–27.09	yes	–	–	–	13.2	64
		seed mean	24.55			–				
DeepSeek Coder 6.7B	16.61	101	14.86	12.81–17.12	yes	13.54	11.86–15.41	no	57.6	810
		102	14.37	12.39–16.49	no	13.28	11.63–15.01	no	58.8	813
		103	14.62	12.52–16.75	yes	13.47	11.72–15.17	no	58.1	816
		seed mean	14.62			13.43				

Paper protocol exactly (caps 2048/64, temperature 0.7/0.01, top- k 20, top- p 0.9), one frozen 200-prompt-per-dataset subset, three independent generation seeds per model, frozen 2024-01-10 registry. Intervals are dataset-stratified prompt-cluster bootstraps. The generic parser is the historical parser for CodeLlama; the DeepSeek family parser is the historical parser for DeepSeek and therefore the acceptance condition. Parser v2 malformed shares are diagnostic; the 64-token cap truncates DeepSeek package queries so often that most of its responses fail the strict grammar. Cap hits are summed over all twelve phases.

Table 4: Track A: seed-level replication estimates against the published values.




Figure 2: Track A replication. The dashed line is the published value; each dot is one seed with its 95% prompt-based confidence interval. The family parser is DeepSeek’s original parser, and every family-parser interval excludes the published value.

out 1.5–2 percentage points low at every seed. That consistent bias points at the model artifacts themselves: the 2024 study used GPTQ 4-bit weights through transformers, while we used Q4_0 GGUF weights through Ollama, with possibly

different chat templates. We now record digests so this can be pinned in future, but the 2024 checkpoints cannot be recovered. The malformed shares in Table 4 also show why the old protocol is hard to reuse. Its 64-token cap on package answers cuts DeepSeek off so often (810–816 cap hits per seed) that most of its answers fail a strict grammar.

4.3 Controlled cap diagnostic

Our earlier report [1] had found that DeepSeek’s measured rate rose by 11.5 percentage points when the token cap was raised, and read that as truncation hiding invented names at the end of long lists. The independent audit [17] found a different explanation. With more tokens, DeepSeek drifted out of the requested list format (code fences appeared in 338 of 800 long answers versus 132 of 800 short ones), and the permissive parser then scored code fragments and prose as invented packages. That finding called for a controlled test. We queried both DeepSeek models over the same code and prompts at temperature 0 with a fixed seed and one worker, under caps of 64, 128, 256, 512, and 2048 tokens in counterbalanced order, and checked that shorter answers were prefixes of longer ones. Parser v2 scored everything (Figure 3; Appendix Table 12).

DeepSeek Coder V2 16B scored between 7.10% and 7.48% at every cap. No cap differs from the 2048-token cap by more than 0.33 percentage points, and every paired interval includes zero. DeepSeek Coder 6.7B does not move in one direction either: 5.28%, 5.81%, 6.72%, 5.12%, and 4.78%. Every paired interval includes zero, although the




Figure 3: Controlled cap diagnostic. (A) Unregistered rate at each cap on package answers. (B) Share of malformed answers at each cap. (C) Each cap compared with the 2048-token answer to the same prompt, with 95% paired, dataset-stratified bootstrap intervals. Every interval includes zero. No equivalence margin was preregistered.

interval at cap 256 is wide (-0.04 to 4.29 percentage points). The malformed share barely moves with the cap (about 12% for one model and 53% for the other). What the data support is this: the large cap effect from the earlier work does not reappear as a consistent or monotone effect when decoding is deterministic and the parser is strict. What the data do not support is a claim of “no effect.” We preregistered no equivalence margin, some paired answers still diverged before the truncation point, and we did not vary temperature along with the cap. Two earlier findings do still stand: a 64-token cap silently produces empty answers from a reasoning model (gpt-oss 20B returned nothing on 15 of 15 probe queries, versus 0 of 15 for a CodeLlama control), and a list cut off at a cap interacts with any parser.

4.4 Modern benchmark

Table 5 and Figure 4 give the six Track B cells. Three models are low: Claude Opus 5 at 1.31% (95% interval 1.04–1.62), GPT-5.2 at 1.88% (1.49–2.30), and gpt-oss 20B at 2.15% (1.76–2.57). DeepSeek Coder V2 16B is high at 10.30% (8.77–11.89), with the most malformed answers (18.25%) and the fewest packages per prompt. Grok 4.6 at 8.22% and DeepSeek V4 Flash at 17.91% have intervals so wide (1.55–18.27 and 5.71–29.86) that the point estimates say little on their own. We do not call Opus, GPT-5.2, and gpt-oss a “frontier trio.” gpt-oss is a local 20B open-weight model, and six hand-picked cells are not a representative sample of anything.

Coverage differs a lot between models. gpt-oss returned 111 empty answers (6.94%) even with a 2048-token cap. Grok and DeepSeek V4 Flash gave malformed answers 10–13% of the time, and DeepSeek Coder V2 18% of the time. Because malformed answers add nothing to the rate, we never publish a rate without these columns.

4.5 Failure regimes and tail risk

The central finding is that similar averages can hide very different risks (Figures 5 and 6). If we rank prompts by how many unregistered names they produced and add up their share, three patterns appear.

Low rate, spread out. For Opus, GPT-5.2, and gpt-oss, the three worst answers supply only 5–11% of unregistered names. Removing them moves the rate by 0.10–0.13 percentage points (1.31 to 1.18, 1.88 to 1.74, 2.15 to 2.05). It takes 29–38 prompts to reach half of each cell’s total. The median prompt has no unregistered name, and the worst prompt has 3–7.

High rate, spread out. DeepSeek Coder V2’s three worst answers supply only 3.98% of its unregistered names. Removing them leaves 9.93%; removing ten leaves 9.43%. Its problem is everywhere: 21.0% of prompts produce at least one unregistered name, the highest prompt risk in the group.

Catastrophic tail. A single Grok 4.6 answer, with 350 names in it, supplies 63.06% of Grok’s unregistered names, and three answers supply 82.88%. Remove those three and Grok drops from 8.22% to 1.61%; remove ten and it is 0.80%. For DeepSeek V4 Flash, two prompts supply half of the total, three answers supply 72.42%, and the rate without them is 5.71%. Both models have the *lowest* prompt risk in the group (5.5% and 18.1%). A typical Grok or V4 Flash answer is fine; the tail is a runaway inventory of names. That same concentration is what makes their intervals so wide.

The practical point is that the same average calls for different defenses. Spread-out errors call for checking every recommended name against the registry at install time. Catastrophic tails call for detecting floods, capping the length of recommendation lists, and refusing to enumerate. A leader-

Model	Rate (%)	95% CI	Risk (%)	Malf. (%)	Empty (%)	Pkgs/prompt	Parsed	Unique	Top-3 (%)	Leave-3 (%)
Claude Opus 5	1.31	1.04–1.62	12.1	3.75	0.31	11.60	9,280	95	10.7	1.18
GPT-5.2	1.88	1.49–2.30	12.1	0.56	0.00	9.03	7,222	102	8.1	1.74
gpt-oss 20B	2.15	1.75–2.58	13.1	0.44	6.94	7.78	6,223	91	5.2	2.05
Grok 4.6	8.22	1.55–18.27	5.5	12.56	0.06	8.44	6,754	532	82.9	1.61
DeepSeek Coder V2 16B	10.30	8.77–11.89	21.0	18.25	0.00	3.05	2,437	196	4.0	9.93
DeepSeek V4 Flash	17.91	5.71–29.86	18.1	10.38	1.00	7.57	6,054	1,045	72.4	5.71

Rate = unregistered-PyPI recommendation occurrences / Parser v2-accepted package occurrences (absent from both the 2024-01-10 and 2026-08-12 snapshots and not a standard-library module). Intervals: 50,000-replicate dataset-stratified prompt-cluster percentile bootstrap. Prompt risk = share of the 800 prompts with at least one unregistered recommendation. Top-3 share = share of unregistered occurrences supplied by the three worst responses; leave-3 = rate after removing them (a concentration diagnostic, not an alternative estimand). $n = 1,600$ package responses per cell.

Table 5: Track B primary outcomes with response coverage and concentration.




Figure 4: Rates and response coverage for the six Campaign 4 cells, the two preregistered cloud extension cells, and the Kimi K3 successor cell. (A) Primary rate with its 95% prompt-based interval. (B) Prompt risk. (C) Share of answers that were lists, malformed, or empty. Dotted lines separate the three families, which were compared under separate Holm corrections.

board that puts Grok 4.6 (8.22%) next to DeepSeek Coder V2 (10.30%) hides that they fail in different ways.

4.6 Pairwise sensitivity

Table 6 lists all 15 paired comparisons. Seven pass under the preregistered rule, and five pass under the recentered-null check. Five differences hold up under both:

- Opus versus gpt-oss (−0.84 percentage points, 95% interval −1.28 to −0.39);
- Opus, GPT-5.2, and gpt-oss, each versus DeepSeek Coder V2 (about −8 to −9 percentage points);
- Opus versus DeepSeek V4 Flash (−16.59 percentage points, −28.65 to −4.43).

GPT-5.2 and gpt-oss versus DeepSeek V4 Flash pass the preregistered rule but not the recentered check (adjusted $p = 0.0596$). Opus versus GPT-5.2 is inconclusive (−0.57 percentage points; raw $p \approx 0.008$, adjusted ≈ 0.07). An earlier draft had called that pair significant, based on a bootstrap with 2,000 replicates that did not stratify by dataset; the reviewer’s challenge was verified and adopted. No comparison involving Grok 4.6 survives correction, because

a handful of prompts drive its estimate. We therefore do not rank the six models from best to worst, and we lead with effect sizes, intervals, and the shape of the tail.

4.7 Concurrent-work bridge

Table 7 shows what happens when our outputs are scored the way Churilov scores them. Under our two-registry, non-stdlib definition, the original six Campaign 4 cells score 8.30% (GPT-5.2), 8.87% (DeepSeek Coder V2), 8.89% (DeepSeek V4 Flash), 9.52% (Opus 5), 10.61% (gpt-oss), and 17.49% (Grok 4.6). The extension cells score 6.51% (Qwen 3.5), 7.38% (Kimi K2.7), and 8.37% (Kimi K3). Only Qwen lands inside Churilov’s reported Python range of 5.49–7.27%; Kimi K2.7 is 0.11 percentage points above its top. The two studies share no model, so we describe this contrast rather than test it.

Four facts shape the interpretation. First, using the same extractor does not by itself bring the studies into agreement: on identical outputs, the original six cells stay above the reported band. Second, registry drift is not the explanation: the 2026 registry and the standard-library rule remove only




Figure 5: Concentration curves. Each curve shows what share of a cell’s unregistered names comes from its k worst prompts. (A) Campaign 4 cells. (B) Extension cells. Grok 4.6 and DeepSeek V4 Flash reach half of their total within one and two prompts; the low-rate cells and DeepSeek Codex V2 need 29–52 prompts.

Pair	Δ (pp)	95% CI	Holm p (tail)	sig.	Holm p (recent.)	sig.
Opus 5 vs. DS Codex V2	-8.98	-10.58—-7.45	0.0003	yes	0.0003	yes
Opus 5 vs. DS V4 Flash	-16.59	-28.65—-4.43	0.0003	yes	0.0453	yes
GPT-5.2 vs. DS Codex V2	-8.42	-9.98—-6.89	0.0003	yes	0.0003	yes
GPT-5.2 vs. DS V4 Flash	-16.02	-28.17—-3.80	0.0003	yes	0.0596	no
gpt-oss vs. DS Codex V2	-8.15	-9.71—-6.65	0.0003	yes	0.0003	yes
gpt-oss vs. DS V4 Flash	-15.75	-27.79—-3.54	0.0003	yes	0.0596	no
Opus 5 vs. gpt-oss	-0.84	-1.28—-0.39	0.0029	yes	0.0026	yes
Opus 5 vs. GPT-5.2	-0.57	-1.00—-0.15	0.0653	no	0.0782	no
Opus 5 vs. Grok 4.6	-6.90	-16.84—-0.24	0.1652	no	0.5776	no
GPT-5.2 vs. Grok 4.6	-6.33	-16.39—-0.30	0.5882	no	0.7363	no
gpt-oss vs. Grok 4.6	-6.06	-16.07—-0.56	0.7592	no	0.7500	no
GPT-5.2 vs. gpt-oss	-0.27	-0.72—-0.19	0.8656	no	0.8402	no
Grok 4.6 vs. DS Codex V2	-2.08	-9.07—-8.02	0.8656	no	0.8402	no
Grok 4.6 vs. DS V4 Flash	-9.69	-24.43—-5.72	0.8656	no	0.8402	no
DS Codex V2 vs. DS V4 Flash	-7.61	-19.41—-4.34	0.8656	no	0.8402	no

Difference = first minus second model in percentage points on the 800 shared prompts; 50,000-replicate dataset-stratified paired prompt-cluster bootstrap. The bootstrap-tail convention is the preregistered one; the recentered-null convention is a labeled post-hoc sensitivity. Holm correction spans all 15 pairs. Values of 0.0003 sit at the $15 \times 1/50,000$ resolution floor and are upper bounds.

Table 6: Track B paired comparisons under both bootstrap conventions.

11 of 632 flagged mentions. Third, imports account for 608 of the 621 remaining flags (97.91%). Explicit install commands are rare in generated code, so the open question is whether an import can be mapped straight to a PyPI package at all. Many real packages import under a different name than they install under, and an import can refer to a file in the same project. Fourth, the prompt source matters a great deal. Across the nine cells, the LLM-written datasets score 3.70–26.45% while the Stack Overflow datasets score 0.00–7.50% (Appendix Table 14). That agrees in direction with Churilov’s observation that synthetic prompts produce higher rates. One caution: Kimi K2.7’s and K3’s package-question rates are a different quantity, dominated by extreme Query 2 lists, and must not be compared with code-import rates at all.

4.8 Prompt-conditioned overlap

Table 8 and Figure 7 separate ordinary overlap from overlap on the same prompt. Among the six Campaign 4 cells, five unregistered names are produced by every model, but four of those five are produced by all six models on at least one identical prompt. Across all nine cells, the average Jaccard similarity between pairs is 0.038, two names appear in every model, and one of the two appears in every model on the same prompt. For most pairs, 60–95% of the names they share occur on a shared prompt; the share drops only where a Kimi flood is involved (Appendix Table 15). Ordinary overlap cannot tell whether two models share a name because they share training data, because a package by that name is a common misconception, or because the same prompt pulled it out of both. The aligned numbers show that the prompt is a live explanation. It has to be ruled out before overlap can

Model	Mentions	Imports (%)	Frozen rate (%)	95% CI	Dual rate (%)	95% CI	Cov. (%)	Risk (%)	Unique
Claude Opus 5	788	95.4	9.52	6.04–13.71	9.52	6.04–13.71	50.5	6.1	50
GPT-5.2	759	89.3	8.30	5.66–11.41	8.30	5.66–11.41	49.2	6.0	44
gpt-oss 20B	782	92.3	10.74	8.34–13.28	10.61	8.22–13.14	52.9	8.5	51
Grok 4.6	766	98.2	18.15	12.45–23.92	17.49	11.91–23.18	48.5	7.0	95
DeepSeek Coder V2 16B	778	95.0	9.13	7.02–11.38	8.87	6.80–11.08	56.0	7.9	48
DeepSeek V4 Flash	742	98.9	8.89	6.73–11.23	8.89	6.73–11.23	52.0	7.4	44
Qwen 3.5 Cloud	169	98.8	7.10	3.57–11.05	6.51	3.12–10.38	14.1	1.4	7
Kimi K2.7 Code	745	96.8	7.38	5.50–9.46	7.38	5.50–9.46	53.2	6.5	36
Kimi K3	777	97.4	8.62	6.57–10.84	8.37	6.33–10.57	52.6	7.2	41

Churilov’s public extraction design (a `pip install regex` plus top-level imports, AST-first with a tree-sitter fallback, `stdlib-list 3.11` exclusion, and the artifact’s false-positive filter) applied to the frozen generated code of every cell; no responses were regenerated. Mentions are the unit of the rate; intervals cluster by prompt and stratify by dataset (50,000 replicates). Across all nine cells, 632 frozen-absence mentions reduce to 621 under the dual-registry/non-stdlib definition (11 removed), and imports supply 97.9% of dual-registry flags. Coverage = share of prompts with any extracted reference. Churilov’s reported Python range is 5.49–7.27%.

Table 7: Code-import bridge: frozen Campaign 4 and extension outputs rescored with Churilov’s extractor.




Figure 6: Leave-top- k : the rate after removing the k worst answers, for k of 0, 1, 3, and 10. Flat lines mean the errors are spread out; steep lines mean a few answers dominate. This is a description of concentration, not a replacement for the primary rate.

Overlap statistic	Six cells	Nine cells
Mean pairwise ordinary name Jaccard	0.054	0.038
Names emitted by every model	5	2
of which emitted by every model on one identical prompt	4	1
same-prompt share of the universal set (%)	80	50

Base set: dual-registry, non-stdlib unregistered recommendations accepted by Parser v2. The six-cell column is the frozen Campaign 4 cohort (artifact at commit eac1434); the nine-cell column adds the three extension cells. Candidate names are withheld.

Table 8: Ordinary versus prompt-aligned cross-model overlap.

be read as evidence about training data.

4.9 Preregistered Qwen/Kimi cloud extension

Table 9 reports the two preregistered cloud cells, and Appendix Table 16 their family of 13 comparisons (10 pass the preregistered rule, 9 the recentered check). Qwen 3.5 Cloud is clean: 2.77% (2.07–3.54), 7.88% prompt risk, 13.25%

Outcome	Qwen 3.5 Cloud	Kimi K2.7 Code
Unregistered recommendation rate (%)	2.77	23.10
95% prompt-cluster interval	2.07–3.54	20.19–26.19
Prompt risk (%)	7.88	43.00
Malformed / empty responses (%)	13.25 / 0.00	14.75 / 0.00
Parsed package occurrences	2,782	111,842
Unregistered occurrences	77	25,838
Unique unregistered names	71	21,631
Cap hits: code / package responses	1 / 0	10 / 323
Query 1 rate (%) / occurrences	0.33 / 302	21.82 / 3,887
Query 2 rate (%) / occurrences	3.06 / 2,480	23.15 / 107,955
Responses with ≤ 10 packages: rate (%)	2.82	4.95
Responses with > 25 packages: rate (%) / occurrences	0.00 / 31	23.93 / 106,980
Response-macro mean / median (%)	3.23 / 0.00	8.36 / 0.00
Prompt-macro mean / median (%)	3.75 / 0.00	11.78 / 0.00
Top-3 response share (%)	13.0	4.6
Cap-proxy-excluded rate (%) [post hoc]	2.77	13.18
Tokens: prompt / completion	731,235 / 92,619	1,026,611 / 1,134,677

Both cells: the Campaign 4 fixed 800-prompt subset, Parser v2, dual registries, caps 4096/2048, `think: false`, 2,400 calls each with zero request errors. The cap proxy is the frozen within-phase longest-response proxy because the first extension runner retained cap counts only at phase level; it is explicitly post hoc.

Table 9: Preregistered Qwen 3.5 / Kimi K2.7 extension outcomes with query volume and cap sensitivity.

malformed answers, no package answer at the token cap, and 2,782 accepted names. Kimi K2.7 Code Cloud is 23.10% (20.19–26.19), with 43.00% prompt risk, 14.75% malformed answers, 323 package answers at the cap, and 111,842 accepted names, of which 107,955 come from Query 2.

The two questions have to be looked at separately. Kimi’s Query 1 rate is 21.82% over 3,887 names and its Query 2 rate is 23.15% over 107,955. Even the Query 1 figure is driven by six answers that run past position 100; leaving out the longest answers in each phase (our stand-in for cap hits) brings it to 3.77%. Answers with at most ten packages score 4.95%. Answers with more than 25 packages supply 106,980 names at 23.93%. The weighting matters: the pooled rate is 23.10%, but the average rate per answer is 8.36%, the




Figure 7: Ordinary versus prompt-aligned overlap among the nine cells. (A) Jaccard similarity of the sets of unregistered names. (B) Share of each pair’s shared names that both models produced on at least one identical prompt. Kimi’s flood-inflated name sets drive its low ordinary Jaccard.

average per prompt is 11.78%, and the median of both is zero. These are descriptions, not replacements for the preregistered estimate.

List position is the clearest mechanism (Figure 8; Appendix Table 18). Kimi’s Query 2 rate is 5.28–5.87% through position 25, then 10.85% at positions 26–50, 15.03% at 51–100, and 27.33% after 100. This is a model that sometimes ignores the instruction and enumerates, and whose lists get worse the longer they run; it is not a flat rate. Parser noise plays a part but does not explain it. The 323 answers we treat as cap hits are malformed 26.94% of the time, versus 11.67% for the rest, and they hold 36.86% of all Kimi’s malformed answers. Yet 236 of them do parse as lists and supply 104,420 names. Even after all 323 are removed, Query 2 still climbs from roughly 5–7% through position 25 to 22.03%, 36.94%, and 41.38% at positions 26–50, 51–100, and 101+. Those late bins rest on only 16, 12, and 10 answers, so they show the mechanism rather than measure it precisely. The 234 Query 2 answers with more than 100 packages are not copies of one list: their mean pairwise Jaccard is 0.066, a core of 39 names recurs, and the rest varies with the prompt. The paired Query 1/Query 2 Jaccard for these floods is only 0.005. Because K2.7’s runner did not record which answers hit the cap, every statement above that depends on the cap is a post-hoc approximation. The K3 cell was instrumented to remove that ambiguity.

4.10 Kimi K3 successor test

The Kimi K3 cell of 2,400 calls passed all twelve checks on its per-answer records: one served model id, no errors or request adjustments, 72 exact cap hits (46 of them in package

answers, all in Query 2), and \$16.27 of token-based spend (Table 10). Its primary rate is 33.30% (27.22–38.81), with 21.50% prompt risk, 4.19% malformed answers, and 25,889 accepted names. All eight comparisons in its family put K3 above the other model on this rate (Appendix Table 17); K2.7 minus K3 is –10.20 percentage points (–16.43 to –3.44).

The two questions could not be more different: Query 1 is 1.76% over 1,303 names and Query 2 is 34.97% over 24,586. K3 answers the question about its own code well, and sometimes treats the open-ended question as an invitation to list everything. The exact cap records explain the headline. The 46 package answers that hit the cap are 2.88% of package answers, but they supply 13,408 names (51.79% of all accepted names) and 6,985 unregistered names (81.02% of the total), at a rate of 52.10% within those answers. Without them the rate would be 13.11%. We keep 33.30% as the primary result. Whether an answer hit the cap is only known after the answer was generated, so removing those answers is conditioning on an outcome [12], and we label it a diagnostic, not a correction.

Compared with K2.7, the ranking flips depending on the metric (Table 11, Figure 9). K3 has lower prompt risk (21.50% versus 43.00%), a lower average rate per answer (4.18% versus 8.36%), and only 54 Query 2 answers above 100 packages versus 234. On the 800 shared prompts, the paired difference in Query 2 prompt risk is 20.63 percentage points (16.50–24.75), and the difference in flood frequency is 22.50 percentage points (19.00–26.00). Yet K3’s rarer floods are much worse (52.10% unregistered, versus 23.81% in K2.7’s approximate cap hits), so its pooled rate is higher. Set aside the cap tails of both, as a diagnostic, and the two rates nearly coincide: 13.11% and 13.18%. Twenty-two prompts




Figure 8: Kimi K2.7’s rate is a story about long lists. (A) Query 2 rate by list position. Kimi stays near 5–6% through position 25 and rises to 27.33% after position 100; Qwen 3.5 produces nothing after position 50. (B) Most of Kimi’s Query 2 names sit past position 100 (log scale). (C) The two longest quartiles of answers carry both the volume and the higher rates. Counts are accepted names, not prompts.

Outcome	Kimi K3
Combined unregistered recommendation rate (%)	33.30
95% prompt-cluster interval	27.22–38.81
Prompt risk (%)	21.50
List coverage / malformed / empty (%)	95.81 / 4.19 / 0.00
Parsed / unregistered occurrences	25,889 / 8,621
Response-macro mean / prompt-macro mean (%)	4.18 / 5.78
Query 1 rate (%) / occurrences	1.76 / 1,303
Query 2 rate (%) / occurrences	34.97 / 24,586
Responses with ≤ 10 packages: rate (%)	2.67
Responses with > 25 packages: rate (%) / occurrences	43.62 / 19,227
Exact cap hits: all phases / package responses	72 / 46
Package cap hits: share of responses / parsed / unregistered (%)	2.88 / 51.79 / 81.02
Rate within package cap hits (%)	52.10
Rate excluding package cap hits (%) [post hoc]	13.11
Top-3 / top-10 response share (%)	14.88 / 40.59
<i>Response-metadata gate</i>	
Ordered sidecar rows / missing	2,400 / 0
Served model ids	kimi-k3 only
Request errors / request adjustments	0 / none
Prompt / completion tokens	1,410,107 / 802,768
Token-derived analytic spend (smoke excluded)	\$16.271841

Design attested as fixed before outcome inspection, but not committed before collection; provenance and completion-gate code changed during collection and before scoring. The prespecified 33.30% estimate remains primary; cap-hit exclusion is post-hoc conditioning on a post-treatment variable.

Table 10: Kimi K3 successor cell: prespecified outcomes and the exact response-metadata audit.

flood in both models, against 15.8 expected if flooding were independent (odds ratio 1.73), so some prompts are prone to floods in both. But the Jaccard overlap of flood prompts is 0.083 against a maximum of 0.231 given the two models’ flood rates, so the prompt alone does not decide which model floods. The 54 K3 floods contain 8,010 distinct unregistered names with a mean pairwise Jaccard of 0.0011, and no

Diagnostic	Kimi K2.7	Kimi K3
<i>Aggregate decomposition (all package responses)</i>		
Occurrence-weighted rate (%)	23.10	33.30
Prompt risk (%)	43.00	21.50
Response-macro mean (%)	8.36	4.18
Responses with ≥ 100 packages	240	54
Package cap hits (K2.7: proxy; K3: exact)	323	46
share of responses / parsed / unregistered (%)	20.2 / 93.4 / 96.2	2.9 / 51.8 / 81.0
Rate inside the cap tail (%)	23.81	52.10
Rate excluding the cap tail (%) [post hoc]	13.18	13.11
<i>Shared-prompt Query 2 diagnostics (K2.7 minus K3, pp; 95% paired CI)</i>		
Prompt risk	41.12% vs. 20.50%	20.62 (16.50–24.75)
Responses over 100 packages	29.25% vs. 6.75%	22.50 (19.00–26.00)
Flood contingency (both / K2.7 only / K3 only / neither)	22 / 212 / 32 / 534	
Expected joint floods under independence; odds ratio	15.8; 1.73	
Flood-prompt Jaccard; maximum given marginals	0.083; 0.231	
<i>K3 flood diversity (54 Query 2 responses above 100 packages)</i>		
Unique normalized packages / unregistered names	12,127 / 8,010	
Mean pairwise Jaccard: package sets / unregistered sets	0.0408 / 0.0011	
Unregistered names in ≥ 25% of floods	0	

K3 cap membership is exact (finish reason and token counts per response); K2.7 cap membership is the frozen within-phase longest-response proxy. Paired intervals resample prompts within each of the four 200-prompt datasets (50,000 replicates). Cap exclusion is diagnostic conditioning, not a causal estimate.

Table 11: Kimi K2.7 versus Kimi K3: aggregation decomposition, shared-prompt flood contingency, and cap-tail influence.

name appears in even a quarter of them. The tail is not one memorized list repeated, and it is not a parser stuck in a loop. That does not remove the parser’s ordinary limitations.

The rise with list position survives the removal of exact cap hits: K3’s Query 2 rate goes from 2.70–5.62% through position 25 to 9.06%, 11.22%, and 36.75% at positions 26–50, 51–100, and 101+. So the mechanism has two stages. Occasionally the model runs away and enumerates, and the further down the list it gets, the less real the names become. The token cap reveals that process and cuts it off; nothing




Figure 9: Kimi K2.7 versus Kimi K3. (A) The ranking depends on the metric: K3 is worse when every name counts equally, but better on prompt risk, on the average rate per answer, and after cap-tail exclusion. (B) K3’s 46 exact package cap hits are 2.9% of its answers but 81.0% of its unregistered names. (C) Rates still rise with list position after removing the K2.7 approximate cap hits and the exact K3 cap hits.

here shows that the cap causes it.

5 Discussion

5.1 The parser is part of the measurement instrument

Strict parsing bought precision by refusing unclear answers, and the cost was visible: 79.75% recall in validation, and malformed shares from 0.44% (gpt-oss) to 18.25% (DeepSeek Coder V2) in the benchmark. Differences that large can bias a comparison between models if they are left out, because the rate of a low-coverage model is computed only from the answers that stayed in format. Control tokens leaked by the transport layer should be stripped before parsing and logged. The Grok `<|eos|>` case shows that a hosted endpoint can move a malformed rate by eleven points without moving the primary rate at all.

5.2 Registry time changes the label space

A name can move from “hallucinated” to registered, or from registered to deleted, with no change in the model, and the 2024 registry contained squatted standard-library names that hid module confusion entirely. Registry snapshots should be dated and hashed like model weights, and rates should be scored against both a frozen and a current snapshot with an explicit standard-library rule. The bridge adds a corollary: with the extractor held fixed, registry drift explains almost none of the gap between the two studies, so a disagreement between two 2026 rates is about which models, which prompts, and how names are extracted, not about the calendar.

5.3 Average rates conceal operationally different risks

Six modern models showed three kinds of failure: low and spread out, high and spread out, and catastrophic tail. Security evaluations need tail-aware metrics because the defenses differ. Spread-out risk is handled by checking names against the registry at install time and by asking the user to confirm installations. Catastrophic tails are handled by flood guards and by bounding how many names a model may recommend. The K2.7/K3 comparison is the clearest example: K3 is the safer model on a typical prompt and the worse model when every emitted name counts equally. Neither number contains the other. A report that carries only one of them answers only one operational question.

5.4 What the cap diagnostic resolves, and what it does not

Under deterministic decoding and a strict parser, the proposed cap mechanism for DeepSeek is weakened: rates are flat or non-monotone across a 32-fold range of caps, and every paired interval includes zero. The diagnostic does not show that temperature caused the earlier association [1], and it does not prove equivalence. The exact K3 cap records then show a different relationship between caps and outcomes: answers that hit the cap can dominate a rate, but cap status is only known after generation, so it cannot separate a runaway list from the boundary that truncates it. A randomized within-prompt cap experiment is still needed.

5.5 Replication as measurement debugging

The useful verdict is not a yes or a no. Whether the transport reproduces the artifact, whether the instrument measures package recommendations, and what current models do

are three separate questions with separate evidence. “The pattern reproduces but the exact number does not” is more informative than “replicated” or “failed to replicate,” because it places the leftover gap in the model artifacts rather than in the phenomenon. And that gap only became visible once the parser and cap confounds were removed.

5.6 What the concurrent replication changes

Churilov establishes breadth and a real, registrable cross-model attack surface. Our bridge shows that using the same extractor is necessary but not sufficient for two rates to be comparable. The models, the prompt source, the mapping from import to package, the refusal policy, and the unit of inference all still matter. Prompt-conditioned overlap should become standard before a set of universal names is read as evidence of shared training data, because four of the five all-model names in our cohort are explained by a shared prompt. Two additions would strengthen that study without any new model runs: the current-registry check it defers to future work, and a recomputation of its pairwise tests that clusters by generation.

6 Threats to validity and limitations

- Python only, and a package-question task rather than general code generation.
- The six benchmark cells, the two extension cells, and the K3 cell were chosen by hand. They are not a representative sample of frontier or local models.
- Kimi K3’s scientific design was frozen, but the provenance and completeness checks were expanded during collection; the finished manifest lacks fields added later, which confirms that the running process used the earlier code. Those edits did not touch prompts, responses, parsing, or outcomes, but “frozen pipeline” would overstate the record. The K3 freeze rests on an attestation and a later commit, not on a tamper-evident commit made before the data existed.
- One fixed prompt subset and one generation per model in Track B; only Track A’s seeds tell us how much a rerun would move.
- Parser v2 found 79.75% of visible recommendations in the validation sample. The Track B audit was done by an AI reviewer at the maintainer’s direction; no human read every record.
- The occurrence rate is computed only over names the parser accepted.
- Hosted aliases and provider infrastructure can change even when the served id is recorded. The DeepSeek Coder V2 run was resumed under mixed settings.
- Some package answers hit the token cap, so the catastrophic-tail totals may be lower bounds.

- Kimi K2.7’s cap counts exist only per phase, so its cap analysis uses the longest answers as a stand-in. K3’s cap status is exact, but removing cap hits is still conditioning on an outcome, and the K2.7/K3 comparison mixes exact labels with a stand-in.
 - Kimi K2.7’s rate is dominated by extreme Query 2 lists. The occurrence rate and prompt risk answer different questions.
 - Unregistered PyPI names include real packages from other ecosystems, not only invented strings.
 - Every tail, control-token, sentinel, recentered-bootstrap, bridge, and cap-conditioning analysis was done after seeing the data and is labeled that way. The bridge cohorts share no model, and generated imports can refer to local files.
 - The Qwen 3.5 and Kimi K2.7 cells were preregistered separately and are not folded back into Campaign 4’s original set of comparisons.
- We also list claims that the evidence does not support:
- That frontier models hallucinate at 1.3–2.2%. Grok 4.6 and DeepSeek V4 Flash contradict it, and gpt-oss is not a hosted frontier model.
 - That the cap effect “vanished.” The diagnostic shows non-reproduction, not the absence of any effect.
 - That every unregistered name is invented.
 - That more than five model pairs are definitively different.
 - That Parser v2 is semantically perfect. Its grammar is reliable, but strict abstention, leaked control tokens, and collapsed “no packages” tokens are measurable limitations.
 - That Churilov’s results and ours agree once extraction is matched.
 - That universal names prove shared training data.
 - That Kimi K2.7 has an ordinary recommendation rate of 23.10%, or that removing cap hits “repairs” it.
 - That Kimi K3 is simply worse than K2.7. The ranking depends on the question.
 - That the token cap causes K3’s tail.
 - That the low overlap of flood prompts makes prompt susceptibility irrelevant.

7 Security, ethics, and responsible artifact release

Publishing a list of unregistered candidate names would hand attackers a target list. Following the original study, this paper and every committed artifact release counts, hashes, manifests, model ids and digests, the parser version, and the registry snapshot identities, but never the names. Raw responses are kept outside version control under a content-addressed index and can be shared with verified researchers under a responsible-access policy. This work measures unsafe recommendations. It does not show that any named package is malicious, and a name registered after a model first recommended it cannot be distinguished here from an

unrelated coincidence.

8 Conclusion

The 2024 pattern comes back, but the exact DeepSeek value does not under its own parser. Modern models fail in distinct ways, spread out or in catastrophic tails, and a single average hides the difference; a successor model can be safer on typical prompts while scoring worse overall. Future evaluations of package hallucination should record the registry and parser versions, report how many answers the parser accepted, build confidence intervals on prompts, and measure how concentrated the errors are. The next experiment these results call for is a within-prompt factorial that varies the token cap and adds a bounded-list instruction, launched through the commit-based freeze gate, to separate runaway enumeration from truncation and to test a practical defense.

Disclosure of AI involvement

The code, analyses, and reports in the underlying repository were produced by Anthropic Claude (Opus 5 and Fable 5) and OpenAI Codex under the interactive direction of the author. The two AI systems also performed the delegated label sign-off, the Track B parser audit, and two adversarial reviews of each other’s work, and every adopted finding was re-verified against raw artifacts as committed code. This manuscript was written by Claude (Fable 5.1) from the committed artifacts and the maintainer’s paper outline, and all tables are generated programmatically from the frozen result files. The maintainer’s preregistration sign-offs and delegations are the human control points. The original study, its pipeline, prompt corpus, and detection code are the work of Spracklen et al.

References

- [1] Jacob Anderson. Replicating package-hallucination measurement through a provider-agnostic HTTP client: A partial replication of Spracklen et al. Technical report, versions 1–4 (2026-08-11 to 2026-08-13), repository document REPLICATION.md, 2026. <https://github.com/bojake/PackageHallucination2026>.
- [2] A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-based improvements for inference with clustered errors. *Review of Economics and Statistics*, 90(3):414–427, 2008.
- [3] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, et al. Evaluating large language models trained on code, 2021. arXiv:2107.03374.
- [4] Aleksandr Churilov. The range shrinks, the threat remains: Re-evaluating LLM package hallucinations on the 2026 frontier-model cohort, 2026. arXiv:2605.17062v3; artifact at <https://github.com/churik5/slopsquatting-replication-2026>.
- [5] Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Chapman & Hall, 1993.
- [6] Sarah Gooding. The rise of slopsquatting: How AI hallucinations are fueling a new class of supply chain attacks. Socket blog, 2025. Published 8 April 2025. <https://socket.dev/blog/slopsquatting-how-ai-hallucinations-are-fueling-a-new-class-of-supply-chain-attacks>.
- [7] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979.
- [8] Stuart H. Hurlbert. Pseudoreplication and the design of ecological field experiments. *Ecological Monographs*, 54(2):187–211, 1984.
- [9] Piergiorgio Ladisa, Henrik Plate, Matias Martinez, and Olivier Barais. SoK: Taxonomy of attacks on open-source software supply chains. In *IEEE Symposium on Security and Privacy*, 2023.
- [10] Bar Lanyado. Can you trust ChatGPT’s package recommendations? Vulcan Cyber blog, 2023. <https://vulcan.io/blog/ai-hallucinations-package-risk>.
- [11] Bar Lanyado. Diving deeper into AI package hallucinations. Lasso Security blog, 2024. <https://www.lasso.security/blog/ai-package-hallucinations>.
- [12] Jacob M. Montgomery, Brendan Nyhan, and Michelle Torres. How conditioning on posttreatment variables can ruin your experiment and what to do about it. *American Journal of Political Science*, 62(3):760–775, 2018.
- [13] Andrew Nesbitt. Slopsquatting – when an LLM hallucinates a non-existent package name, and a bad actor registers it maliciously. The AI brother of typosquatting. Credit to Seth Larson for the name. Mastodon post, 2025. Posted 8 April 2025. <https://mastodon.social/@andrewnez/114302875075999244>.
- [14] Shradha Neupane, Grant Holmes, Elizabeth Wyss, Drew Davidson, and Lorenzo De Carli. Beyond typosquatting: An in-depth look at package confusion. In *Proceedings of the 32nd USENIX Security Symposium*, 2023.
- [15] Brian A. Nosek, Charles R. Ebersole, Alexander C. DeHaven, and David T. Mellor. The preregistration revolution. *Proceedings of the National Academy of Sciences*, 115(11):2600–2606, 2018.

- [16] Marc Ohm, Henrik Plate, Arnold Sykosch, and Michael Meier. Backstabber’s knife collection: A review of open source software supply chain attacks. In *Detection of Intrusions and Malware, and Vulnerability Assessment (DIMVA)*, 2020.
- [17] OpenAI Codex. Independent audit of the package-hallucination replication and extension. Repository document `codex-experiment.md` (2026-08-12, with a forensic addendum and a 2026-08-25 update), with the maintainers’ verification in `AUDIT_RESPONSE.md`, 2026. <https://github.com/bojake/PackageHallucination2026>.
- [18] Joseph Spracklen, Raveen Wijewickrama, A H M Nazmus Sakib, Anindya Maiti, Bimal Viswanath, and Murtuza Jadliwala. We have a package for you! A comprehensive analysis of package hallucinations by code generating LLMs. In *Proceedings of the 34th USENIX Security Symposium*, 2025. arXiv:2406.10279v3.
- [19] Donald Stufft. PEP 503 – simple repository API. Python Enhancement Proposals, 2015. <https://peps.python.org/pep-0503/>.
- [20] Markus Zimmermann, Cristian-Alexandru Staicu, Cam Tenny, and Michael Pradel. Small world with high risks: A study of security threats in the npm ecosystem. In *Proceedings of the 28th USENIX Security Symposium*, 2019.

A Artifact checklist

The repository accompanying this paper contains:

- the full preregistrations and their dated amendments (Experiments/PREREGISTRATION_v4.md, PREREGISTRATION_OLLAMA_CLOUD_EXTENSION.md, PREREGISTRATION_KIMI_K3_EXTENSION.md);
- the frozen Parser v2 grammar and its unit tests (parser_v2.py);
- the initial and Track B audit designs, signed summaries, and hashes;
- the construction, dates, and counts of both registry snapshots;
- all seed-level Track A tables and the full cap pairing and position-band tables (campaign4_results.json, cap_diagnostic_summary.json);
- all 15 paired comparisons under both conventions (campaign4_post_analysis.json);
- prompt-level derived counts sufficient to reproduce every interval without exposing raw names;
- the DeepSeek Coder V2 run-history addendum;
- the Churilov-compatible extractor port and its machine-readable output (churilov_bridge_analysis.py, .json);
- the separately frozen extension protocol, manifests, and post-analysis (ollama_cloud_extension_results.json, ollama_cloud_extension_post_analysis.json);
- the Kimi K3 per-answer metadata records, cost ledger, scoring gate, and results (kimi_k3_extension_results.json);
- the K2.7/K3 paired post-analysis (kimi_k2_7_vs_k3_post_analysis.json);
- the retrospective provenance stamp, with its explicit statement that it is not a cryptographic preregistration, and the raw-artifact hash index;
- the review amendment ledger; and
- the commit-based freeze gate (experiment_gate.py) required for every future analytic run.

The manuscript’s tables and figures are regenerated from these files by paper/build_tables.py and paper/make_figures.py.

B Additional cap-diagnostic tables

Table 12 lists every paired cell of the controlled cap diagnostic, and Table 13 the rates by list position within answers that parsed as lists.

Model	Cap	Packages	Unreg.	Rate (%)	Malformed (%)	Δ vs. 2048 (pp)	95% CI
DeepSeek Coder V2 16B	64	323	23	7.12	12.5	-0.02	-0.20–0.11
	128	320	23	7.19	12.0	0.04	-0.11–0.21
	256	324	23	7.10	12.0	-0.04	-1.19–1.06
	512	321	24	7.48	12.5	0.33	-0.08–1.05
	2048	322	23	7.14	12.5	ref.	
Spearman ρ (rate vs. cap)				0.30			
DeepSeek Coder 6.7B	64	265	14	5.28	52.5	0.50	-0.23–1.35
	128	258	15	5.81	52.5	1.03	-0.25–2.66
	256	238	16	6.72	53.0	1.94	-0.04–4.29
	512	254	13	5.12	53.0	0.34	-1.94–2.83
	2048	251	12	4.78	53.5	ref.	
Spearman ρ (rate vs. cap)				-0.60			

100 prompts (25 per dataset), both package queries, temperature 0, seed 7, one worker, counterbalanced cap order, Parser v2, frozen 2024 registry. Deltas pair each cap with the 2048-token response at the prompt level (20,000-replicate dataset-stratified bootstrap). No equivalence margin was preregistered, so a small delta is not proof of no effect.

Table 12: Controlled cap diagnostic: paired cells.

Model	Cap	Query	Pos. 1	Pos. 2	Pos. 3	Pos. 4	Pos. 5+
DeepSeek Coder V2 16B	64	Q1	6.5% (n=92)	3.6% (n=55)	10.0% (n=10)	33.3% (n=3)	0.0% (n=1)
	64	Q2	7.3% (n=82)	9.1% (n=66)	10.0% (n=10)	0.0% (n=4)	–
	128	Q1	6.5% (n=92)	3.7% (n=54)	10.0% (n=10)	33.3% (n=3)	0.0% (n=1)
	128	Q2	7.2% (n=83)	9.1% (n=66)	12.5% (n=8)	0.0% (n=3)	–
	256	Q1	6.5% (n=92)	3.6% (n=55)	10.0% (n=10)	33.3% (n=3)	0.0% (n=1)
	256	Q2	7.2% (n=83)	9.1% (n=66)	0.0% (n=9)	0.0% (n=4)	100.0% (n=1)
	512	Q1	7.6% (n=92)	3.6% (n=55)	10.0% (n=10)	33.3% (n=3)	0.0% (n=1)
	512	Q2	7.3% (n=82)	9.2% (n=65)	11.1% (n=9)	0.0% (n=4)	–
	2048	Q1	6.5% (n=92)	3.7% (n=54)	10.0% (n=10)	33.3% (n=3)	0.0% (n=1)
	2048	Q2	7.2% (n=83)	9.1% (n=66)	11.1% (n=9)	0.0% (n=4)	–
	64	Q1	6.2% (n=64)	0.0% (n=39)	18.8% (n=16)	12.5% (n=8)	16.7% (n=6)
	64	Q2	6.7% (n=30)	0.0% (n=26)	6.2% (n=16)	9.1% (n=11)	2.0% (n=49)
DeepSeek Coder 6.7B	128	Q1	7.8% (n=64)	0.0% (n=39)	18.8% (n=16)	12.5% (n=8)	16.7% (n=6)
	128	Q2	6.9% (n=29)	0.0% (n=25)	0.0% (n=15)	10.0% (n=10)	4.3% (n=46)
	256	Q1	7.9% (n=63)	0.0% (n=38)	20.0% (n=15)	14.3% (n=7)	0.0% (n=4)
	256	Q2	7.1% (n=28)	0.0% (n=24)	0.0% (n=14)	10.0% (n=10)	11.4% (n=35)
	512	Q1	4.8% (n=63)	0.0% (n=39)	18.8% (n=16)	12.5% (n=8)	33.3% (n=3)
	512	Q2	6.9% (n=29)	0.0% (n=25)	0.0% (n=15)	0.0% (n=9)	6.4% (n=47)
	2048	Q1	6.2% (n=64)	0.0% (n=39)	18.8% (n=16)	12.5% (n=8)	16.7% (n=6)
	2048	Q2	7.4% (n=27)	0.0% (n=23)	0.0% (n=13)	11.1% (n=9)	0.0% (n=46)

Unregistered (frozen-registry) rate at each package position within grammar-valid responses, with the number of responses contributing at that position. Q1 = packages required by the generated code; Q2 = packages useful for the task.

Table 13: Controlled cap diagnostic: position-band rates.

C Bridge and overlap detail

Table 14 breaks the code-import bridge down by prompt dataset, and Table 15 gives the fifteen pairwise overlap statistics for the Campaign 4 cells.

Model	LLM-LY	LLM-AT	SO-LY	SO-AT
Claude Opus 5	9.72 (288)	13.83 (311)	2.26 (133)	1.79 (56)
GPT-5.2	11.33 (309)	8.25 (291)	1.92 (104)	3.64 (55)
gpt-oss 20B	11.91 (319)	13.40 (291)	2.83 (106)	4.54 (66)
Grok 4.6	13.29 (301)	26.45 (344)	1.21 (83)	5.26 (38)
DeepSeek Coder V2 16B	9.64 (280)	10.54 (275)	4.89 (143)	7.50 (80)
DeepSeek V4 Flash	10.89 (303)	10.69 (262)	2.50 (120)	3.51 (57)
Qwen 3.5 Cloud	12.70 (63)	3.70 (81)	0.00 (12)	0.00 (13)
Kimi K2.7 Code	8.36 (287)	9.56 (251)	2.24 (134)	5.48 (73)
Kimi K3	10.20 (294)	10.38 (289)	1.61 (124)	4.29 (70)

Dual-registry, non-stdlib rate in percent with the mention count in parentheses. LLM = LLM-synthesized prompts; SO = Stack Overflow prompts; LY/AT = last-year/all-time.

Table 14: Code-import bridge rates by prompt dataset.

Pair	Shared	Name Jaccard	Same-prompt	Share (%)	Aligned Jaccard
Opus 5 vs. GPT-5.2	21	0.119	17	81	0.119
Opus 5 vs. gpt-oss	16	0.094	12	75	0.095
Opus 5 vs. Grok 4.6	12	0.020	9	75	0.026
Opus 5 vs. DS Coder V2	14	0.051	10	71	0.032
Opus 5 vs. DS V4 Flash	23	0.021	19	83	0.021
GPT-5.2 vs. gpt-oss	27	0.163	23	85	0.140
GPT-5.2 vs. Grok 4.6	23	0.038	19	83	0.039
GPT-5.2 vs. DS Coder V2	22	0.080	21	95	0.066
GPT-5.2 vs. DS V4 Flash	38	0.034	36	95	0.036
gpt-oss vs. Grok 4.6	18	0.030	15	83	0.036
gpt-oss vs. DS Coder V2	18	0.067	17	94	0.058
gpt-oss vs. DS V4 Flash	35	0.032	31	89	0.033
Grok 4.6 vs. DS Coder V2	13	0.018	10	77	0.014
Grok 4.6 vs. DS V4 Flash	23	0.015	15	65	0.013
DS Coder V2 vs. DS V4 Flash	39	0.032	34	87	0.028

Shared = unique unregistered names emitted by both models; same-prompt = shared names emitted by both on at least one identical prompt; aligned Jaccard is computed over (normalized name, dataset, prompt index) tuples.

Table 15: Pairwise overlap among the six Campaign 4 cells.

D Extension-cell detail

Tables 16 and 17 give the two extension comparison families, and Table 18 the Query 2 rates by list position with and without cap-tail exclusion.

Pair	Δ (pp)	95% CI	Holm p (tail)	sig.	Holm p (recent.)	sig.
Opus 5 vs. Qwen 3.5	-1.45	-2.19—0.77	≤ 0.00026	yes	0.0006	yes
Opus 5 vs. Kimi K2.7	-21.79	-24.85—18.88	≤ 0.00026	yes	≤ 0.00026	yes
GPT-5.2 vs. Kimi K2.7	-21.22	-24.27—18.31	≤ 0.00026	yes	≤ 0.00026	yes
gpt-oss vs. Kimi K2.7	-20.95	-24.01—18.01	≤ 0.00026	yes	≤ 0.00026	yes
DS Coder V2 vs. Qwen 3.5	7.53	5.93—9.17	≤ 0.00026	yes	≤ 0.00026	yes
DS Coder V2 vs. Kimi K2.7	-12.80	-16.03—9.63	≤ 0.00026	yes	≤ 0.00026	yes
DS V4 Flash vs. Qwen 3.5	15.14	2.96—27.11	≤ 0.00026	yes	0.0415	yes
Qwen 3.5 vs. Kimi K2.7	-20.33	-23.46—17.40	≤ 0.00026	yes	≤ 0.00026	yes
Grok 4.6 vs. Kimi K2.7	-14.88	-22.75—4.19	0.0436	yes	0.0294	yes
GPT-5.2 vs. Qwen 3.5	-0.88	-1.62—0.19	0.0445	yes	0.0582	no
gpt-oss vs. Qwen 3.5	-0.61	-1.37—0.10	0.2828	no	0.3110	no
Grok 4.6 vs. Qwen 3.5	5.45	-1.34—15.49	0.5984	no	0.4888	no
DS V4 Flash vs. Kimi K2.7	-5.20	-17.42—7.19	0.5984	no	0.4888	no

Thirteen comparisons involving a new cell form their own Holm family; Campaign 4’s original 15-pair family is untouched. Values at the $13 \times 1/50,000 = 0.00026$ floor are bounds.

Table 16: Paired comparisons involving the Qwen 3.5 and Kimi K2.7 extension cells.

Pair	Δ (pp)	95% CI	Holm p (tail)	Holm p (recent.)
Opus 5 vs. Kimi K3	-31.99	-37.40—25.93	≤ 0.00016	≤ 0.00016
GPT-5.2 vs. Kimi K3	-31.42	-36.96—25.32	≤ 0.00016	≤ 0.00016
gpt-oss vs. Kimi K3	-31.15	-36.65—25.06	≤ 0.00016	≤ 0.00016
Grok 4.6 vs. Kimi K3	-25.08	-34.68—13.29	0.0004	0.0002
DS Coder V2 vs. Kimi K3	-23.00	-28.68—16.70	≤ 0.00016	≤ 0.00016
DS V4 Flash vs. Kimi K3	-15.39	-28.51—2.00	0.0252	0.0207
Qwen 3.5 vs. Kimi K3	-30.53	-36.11—24.35	≤ 0.00016	≤ 0.00016
Kimi K2.7 vs. Kimi K3	-10.20	-16.42—3.44	0.0068	0.0042

Comparator minus Kimi K3 in percentage points over all 800 shared prompts; separate eight-comparison Holm family. Values at the $8 \times 1/50,000 = 0.000160$ resolution floor are reported as bounds, never as exact p -values.

Table 17: Paired comparisons involving Kimi K3 (all eight survive both conventions).

Position	Qwen 3.5	Kimi K2.7	K2.7, proxy excluded	Kimi K3	K3, exact caps excluded
1	3.40 (735)	5.28 (682)	4.87 (452)	2.57 (779)	2.70
2	4.09 (685)	5.61 (677)	5.59 (447)	4.59 (762)	4.56
3-5	2.46 (936)	5.29 (1,891)	5.08 (1,201)	3.13 (2,044)	2.96
6-10	0.00 (100)	5.65 (2,424)	5.41 (1,275)	3.78 (2,114)	3.65
11-25	0.00 (18)	5.87 (3,882)	6.71 (447)	6.25 (1,360)	5.62
26-50	0.00 (6)	10.85 (6,045)	22.03 (345)	7.72 (1,554)	9.06
51-100	—	15.03 (11,760)	36.94 (509)	15.90 (2,900)	11.22
101+	—	27.33 (80,594)	41.38 (1,073)	59.15 (13,073)	36.74

Occurrence-weighted unregistered rate in percent with parsed occurrences in parentheses, by package position within Query 2 responses. Excluded columns remove the K2.7 longest-response cap proxy or the exact K3 package cap hits; late bins after exclusion rest on few responses and are mechanism evidence, not stable population rates.

Table 18: Query 2 position profiles for the extension cells.